
I don't know how to start this shit, yo.. now

Summary

Write your summary here...

Preface

Write your preface here...

Contents

Summary	i
Preface	ii
Table of Contents	iv
List of Tables	v
List of Figures	viii
Abbreviations	ix
1 Introduction	1
2 Background	5
2.1 Complex Systems	6
2.1.1 Cellular Automata	9
2.2 Material Computing	11
2.3 Computing Neural Networks	14
2.4 Reservoir Computing	15
2.4.1 Linear and nonlinear output layers	17
3 Making Of A Cyborg	19
3.1 Concept	19
3.2 Architecture	20
3.3 Wetware	20
3.3.1 Neural Interface	21
3.4 Software	22
3.4.1 MEAME	22
3.4.2 SHODAN	23
3.4.3 Storage And User Interface	24
3.5 Current State	24

4	Experimental System and Setup	33
4.1	Data Loop	34
4.2	Core Reservoir Computer	34
4.3	Reconfiguration Loop	35
4.4	Parameter Space	37
5	Conclusion And Further Work	39
5.1	Further Work	39
Bibliography		39
Appendix		43

List of Tables

List of Figures

1.1	From socrates webpage	2
2.1	Lorenz attractor. Image taken from Strogatz - Nonlinear dynamics and chaos	8
2.2	From Bar Yam book.	9
2.3	From Sipper paper.	10
2.4	From Bar Yam book.	10
2.5	Source: A new kind of science.	11
2.6	Source: wikimedia commons	11
2.7	Should add a time axis to drive the point home that these are 1D	12
2.8	Langtons egg	12
2.9	A feed forward ANN.	16
2.10	A feed forward ANN.	16
2.11	A feed forward ANN.	17
2.12	Should say “kernel function” or similar on top of the arrow	17
3.1	A simple conceptual cyborg.	20
3.2	An overview of all software written for the cyborg project. The dataflow is color coded with red denoting data going from neurons to robot, green from robot sensors to neurons. The ble line indicates configuration data, as the system is capable of reconfiguring itself during an experiment. The two main systems are physically separated, connected only through network.	26
3.3	Approximated Kolmogorov complexity of the firing patterns over a one year period.	27
3.4	A look through the microscope, showing a culture nicknamed “Frank”, a play on frankensteins monster due to it’s tendency to develop <i>Organoids</i> , small proto-organs with structures similar to eyes and other sensor organs. The black lines are electrodes that can be used to record and apply stimuli.	28
3.5	A placeholder for an open headstage showing the electrodes.	29
3.6	Rough sketch. nice MEAME TODO: Flesh out the DSP box with RPC handler a la MCS USB .dll	29

3.7	The interleaved datastream can be broken down into its individual channels. Here a datastream with only 3 streams are shown, whereas the stream is normally comprised of data from 60 channels.	30
3.8	Rough sketch. Some wavez	30
3.9	Rough sketch. Some wavez	31
3.10	Rough sketch.	32
4.1	Rough sketch. Some wavez	33
4.2	Rough sketch. Some wavez	34
4.3	The cyborg doing its thing	35
4.4	A simple conceptual cyborg.	36
4.5	Rough sketch. Some wavez	36

Abbreviations

ANN = *Artificial* Neural Network
RC = Reservoir Computing

Chapter

1

Introduction

They're trying to understand what space is. That's tough for them. They break distances down into concentrations of chemicals. For them, space is a range of taste intensities.

Greg Bear, Blood Music

To investigate the computational properties of biological neural networks a hybrid biological digital robot has been created in a joint multidisciplinary effort. Countless man-hours have been spent improving the design and manufacturing process of the digital computer, creating more and more complex architectures capable of operating at ever greater speeds. The brain is not subject to this top down design philosophy, yet through a process of self organization neurons are capable of forming highly complex networks capable of solving complex tasks, with far greater energy efficiency, robustness and parallelism than any designed processor. Inspired by work done in the field of material computing systems such as the mecobo platform of nascence [citation needed], a similar system has implemented using living neural networks grown from human stem cells. Neurons are grown *in vitro*¹ on *Micro Electrode Arrays* 1.1:A, C (MEA) are interfaced with a digital computer via voltage spike measurements 1.1:B, forming a hybrid neuro-digital system. This system utilizes the theoretical framework of *Reservoir Computing* to help translate between the digital and biological parts of the system, allowing it to solve simple tasks as a proof of concept. The ultimate goal of the cyborg is to further our understanding of the underlying principles that governing how nature computes.

Complexity

In the 50's and 60's there was much optimism in the burgeoning field of artificial intelligence. In 1965 H. A. Simon claimed "machines will be capable, within twenty years,

¹From latin, literally "in glass", as opposed to in the body.

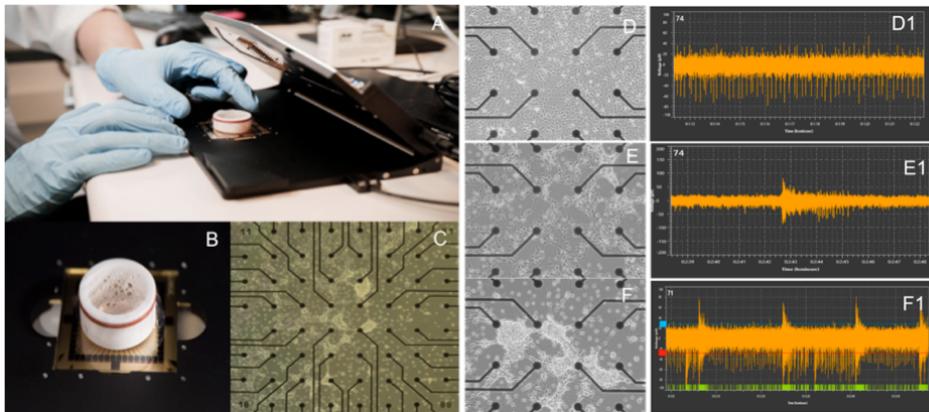


Figure 1.1: From socrates webpage

of doing any work a man can do.”[4] , while Marvin Minsky boldly claimed in 1967 that “Within a generation ... the problem of creating ‘artificial intelligence’ will substantially be solved.” [1]. Had they chosen to predict any other field, such as logistics, information sharing or communications their statements would have been prophetic and visionary, so why did artificial intelligence turn out so differently? The researchers sought to make machines that could use logic similar to that of high level human thinking. Therefore, it followed that the machine had to be programmed with rules governing logic in order to reach sound conclusions. To represent the prior and deduced knowledge, the researchers designed programming languages such as lisp that could accurately describe these operations. In order to actually execute these lisp programs hardware had to be created, supporting the primitive operations such as addition, subtraction and loading from memory. Regardless of the underlying platform a lisp program did not change meaning, separating execution and intent in an elegant manner allowing the programmer to ignore implementation details. In short, each piece of the puzzle was self contained, allowing the researchers to build large systems where each “layer” interacted in a manner that could be reasoned about independently, keeping the *complexity* of the system in check. Nature, on the other hand, applies a completely different method. Complex structures appear with no blueprint, arising from a process of *self-organization* driven by a set of growth rules. This self organizing process is capable of producing incredibly complex, robust and diverse structures whose functionality arises not from specialized components working in isolation, but from the interaction of many components. The early AI researchers believed that to create an intelligent system it was sufficient to provide a platform that could reason logically, but as history shows this did not work out. The human brain is not hiding some underlying logic beneath incidental complexity, instead the complexity seems to be the fundamental driving force behind intelligence. In short, intelligence does not arise from logic, instead our understanding logic arises from our intelligence, and to study intelligence it is necessary to focus on the complex behavior of neurons from which it originates.

Computation

The invention of the digital computer will be remembered as one of, if not the most significant technological advances of mankind. With this neatly designed clockwork machine the concept of computation seemed rather straight forward, where each executed machine instruction corresponded to some unit of computation. This idea of computation is very convenient, and it fails utterly when applied to biological systems. In a recent example of this from 2005 [Cite The Singularity Is Near: When Humans Transcend Biology] Ray Kurzweil claimed that computers would be as powerful as a mouse brain by 2020, a prospect that looks rather unlikely in 2018. Other than in the head of computer scientists, neurons have no notion of floating point operations or branching logic, and measuring the computational capabilities of a brain in FLOPS is grossly underestimating what computing *can* be. In spite of the digital computers shortcoming compared to the human brain, the conventional digital approach to computing has been hugely successful in solving problems that humans are bad at, and is so ubiquitous that other approaches have been dubbed *Unconventional Computing*. Unconventional computing, as implied by the name, comes in many forms such as buckets of water [2], or blobs of carbon nanotubes [cite nascence]. In these unconventional approaches it becomes harder and harder to pin down exactly what computation is and what distinguishes it from any ordinary physical process.

Cyborg

The main body of work done for this thesis is the design and development of a *Cyborg*² capable of sensing and maneuvering a simple maze. This work has been performed as part of the NTNU cyborg project [cite webpage] which comprises of researchers from the department of neuroscience, nano science, computer science, cybernetics and more. More specifically, the contribution to this cross-disciplinary effort presented in the thesis is the implementation of software for a closed loop system in which a robot is controlled by the neural tissue whose sensors feed data back to the neural tissue. Both by necessity and by choice most biological details are left out, the focal point of the thesis is the system in which neural tissue is a part, not the neurons themselves. Thus only details that are necessary for the computational model of neural tissue will be considered, leaving the intricate complexities of neurons to the neuroscientists. TODO: mer her?

²A portmanteau word stemming from cybernetic organism

Chapter 2

Background

The genetic code does not, and cannot, specify the nature and position of every capillary in the body or every neuron in the brain. What it can do is describe the underlying fractal pattern which creates them.

Academician Prokhor Zakharov
Sid Meier's Alpha Centauri

Creating a cyborg is a massive cross disciplinary effort, and if a scope is not clearly defined the background runs the risk of becoming equally massive. The goal of the thesis is to create a hybrid neuro-digital system capable of controlling a simple robot, and by extension to create a “bridge” between neural and digital. Each section in the background shares this bridge as a red thread, and consequently topics that are not directly related to the thesis’ goal are discarded. The background is laid out as follows: First *Complex Systems* are introduced as a framework to discuss the computational capabilities in a wide range of systems that exhibit system dynamics similar to that of neurons. Modeling neurons as a complex systems is a first step towards establishing a common “language” between neural activity and digital logic. Next, *Material Computing and Evolution In Materio*, EiM for short, introduces computation done in unstructured matter through the process of evolution. The goals of EiM are closely aligned to the goal of this thesis, as both study massively parallel computation happening in physical matter shaped by the process of evolution. Next section, *Neurons As Computers* introduces neurons with a strong focus on their computational capabilities. Building on the EiM section, this section motivates a necessary reduction of scope, proposing a simplified model of the computing neuron and a medium of communication between neuron and digital. Finally, *Reservoir Computing* is introduced, tying together the previous sections by introducing the framework of reservoir computing in order to establish a common “language” between neural cultures and digital.

2.1 Complex Systems

Before discussing complex systems it is necessary to provide a definition of complexity and system. In common parlance complexity one might describe a book as complex due to its intertwined narrative, where small details later turn into large plot points. A song can be described as complex because of how the different instruments play together, providing a context to the lyrics and vice versa. Giving an exact of what complexity is is awkward, but the common theme is highly interconnected interaction. In the first example this refers to how the narratives intertwine, where no narrative makes sense outside of the context of the other narratives, while in the second example it refers to how the different instruments can interplay, giving meaning to each other in such a way that one instrument alone, or the lyrics simply written down loses meaning and impact.

A system is defined as a set of elements and the relations between them. When a system changes over time it is a *dynamical* system, a property that is implicitly assumed for the rest of this thesis. When there are no interaction the system can only be described by its constituent elements, and these systems are not very interesting. When systems have interaction this interaction can be either linear or *nonlinear*. Linear systems are characterized by the independence of interaction. In these systems the behavior of the system as a whole is driven by the interaction of components, but each individual component behaves independantly of the rest of the system. When designing systems this is a desirable property because it enables a *top down* design process where each part of a machine can be designed independently before being put together, as each part does not alter its behavior when part of a larger system, allowing us to assemble large systems that are *complicated* while keeping complexity low. When the interactions become nonlinear this property is lost, and these systems become very hard to reason about and design. As Strogatz puts it, “Whenever parts of a system interfere, cooperate or compete there are nonlinear interactions going on. ...if you listen to your two favorite songs at the same time you won’t get double the pleasure!” [cite Strogatz, page 9]. In *complex systems* nonlinear interaction is what shapes the systems behavior which becomes greater than the sum of parts.

Feedback and Self Organization

The nonlinearity of interactions in complex systems is often expressed as *feedback* which when positive can amplify small perturbations into cascading effects that change the system entirely, or when negative act as dampeners, givin the system some measure of stability. An example of this is how bees harvest nectar which must be carefully dried, a process that can be done on a large scale if all the nectar has the same water content. Because of this a hive will harvest from only one source at a time, which in beekeeping terms is known as “flow” (with raspberry-flow being the most important one). With no centralized decisionmaker the bees are still able to quickly agree on a new flow once the current flow dries up through local interactions. When a worker bee finds a suitable source of nectar it will fly back to the hive and perform a dance, informing other bees of the food source. Other bees that see this dance will then seek out the food source, and if successful they too will return to the hive to recruit new workers, however even after a successful foraging they can still be recruited to a different source. Often several food sources are found, and these sources “compete” against each other, but typically the closest source wins because the

faster trips cause recruitment to happen quicker. Through the amplification of local effects (one bee finding a food source) the bees settle on one type of nectar, a *global behavior* that has *emerged* in a *bottom up* fashion, i.e without the need for a centralized decisionmaker through the process of *self organization*. Through collective behavior bees can adapt to changes in their environment, in fact a hive may even survive the death of their queen by moving a worker larvae to a queen cell. To highlight how much of the bees sophistication arise from collective behavior, consider the fact that if a hive is moved even a meter, the forager bees will not be able to find their way back home, clearly as an individual the bee is not a great thinker!

Just like a bee does not operate in a vacuum, the beehive is part of a larger ecosystem. Local interactions on the *scale* of a single bee shapes the behavior at the scale of the entire hive, which in turn interacts with the surrounding flora at the scale of the eco-system. In short, every layer that is peeled away, from the scale of the eco-system to the scale of local flora down to the scale of a single bee reveals a new layer of complexity. Changing the scale of observation does not only reveal new complexities, it can just as well reveal simplicity and order. An example of this *emergent simplicity* is the earth itself which behaves as a simple celestial body in the solar system regardless of the complexities found in smaller scales of observation.

State Spaces and Attractors

The space of possible states, or configurations a system can be in is known as a *state space*¹. As an example, the state space of a rubiks cube can be defined as all possible configurations a cube can be in. This space includes the solved cube as well as any permutation reachable from the solved cube when twisting it, but it does not include impossible rubiks cubes such as a cube with ten red faces. A single system may be defined by multiple state spaces depending on which parts of the system is observed and at which scale. The state space of the rubiks cube can be extended to include the current rotation the cube has in space, or it can be observed at a finer scale where two solved cubes occupy different states if the face of one cube has been slightly rotated. These states are known as *macro states* and have been chosen specifically by the observer as opposed to the *micro state* which is the ground truth of the system. For physical systems the micro state is not only intractable, but impossible to measure due to the uncertainty principle, and generally what happens at a subatomic level is not very relevant when describing a rubiks cube.

As a dynamic system evolves over time it follows a trajectory through its state space which depends on its initial conditions and outside perturbation. Analogous to gravity wells in cosmology, certain points in the state space act as *attractors*, “pulling in” the states in its *attractor basin*. The simplest such attractor is the point attractor in which a system in the point attractors basin will move towards a static state. Cyclic attractors are a sequence of states that repeat themselves sequentially. A system in a cyclic attractor can oscillate between two states, or it can follow a longer path, known as cycle length before repeating itself. For both cyclic and point attractors the amount of states a system goes through before reaching the attractor (that is, the amount of states visited only once) is known as the *transient*. The length of the transient is a property of one trajectory leading

¹For continuous systems the word phase space is more accurate, but it's not a useful dichotomy for this thesis.

to the attractor, but the average trajectory length is a property of the attractor itself. There is a third type of attractor which have have no cycles, known as *strange attractors*. Systems in strange attractors exhibits *chaotic* behavior, characterized by its sensitivity to initial conditions and unpredictable, seemingly random behavior. Fig 2.1 shows one solution for a nonlinear differential equation known as Lorenz attractor. Two near-identical initial conditions in a this attractor will quickly diverge in behavior, however both will create the “butterfly” pattern. For a discrete system a strange attractor is impossible as any discrete system can only represent a finite amount of states, however a cyclic attractor with a cycle length near the total amount of expressible states may be informally considered as strange.

TODO: Something about attractor landscape wrt complex, chaotic and ordered systems.

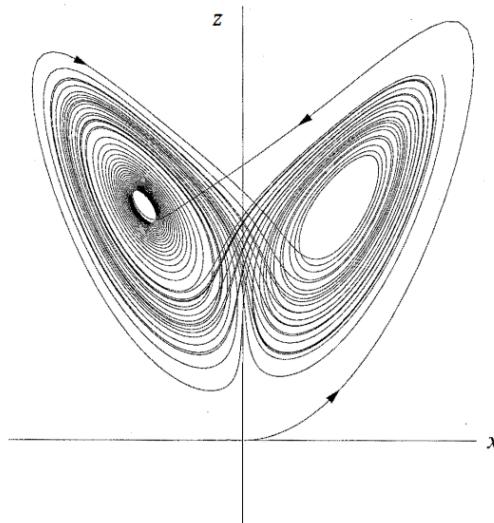


Figure 2.1: Lorenz attractor. Image taken from Strogatz - Nonlinear dynamics and chaos

Universality Of Complex Systems

Not all complex systems are biological in nature. In Dynamics of Complex Systems Yaneer Bar-Yam illustrates this as shown in fig 2.4. In the classical school of thought different fields are seen as divergent, however Bar-Yam points out that each field deals with systems, and it turns out that complex systems have much in common regardless of how and where they originate. This *universality* is a powerful tool when studying systems such as neural networks because it allows us to explain and study much of the behavior with far simpler models, one such model being *cellular automaton*s.

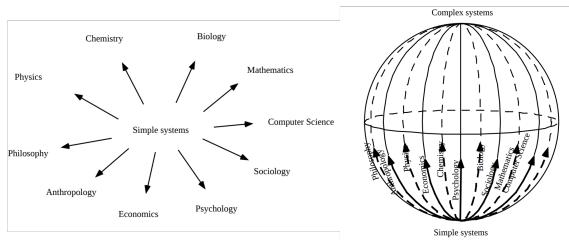


Figure 2.2: From Bar Yam book.

2.1.1 Cellular Automata

A *Cellular Automaton*², CA in short, is a simple discrete model of a single cell which changes between a discrete set of states based only on its immediate neighbors. Figure 2.7 shows an example of a *rule set*, or *transition table* for a cellular automaton with only two states, and an initial configuration, or “seed”, which over several steps forms a pattern. CAs have many properties that make them ideal as a model for studying complex systems and how they can facilitate computation. The reasons for this can be seen in ?? from [cite Sipper] showing how CAs model computation made by simple elements interacting locally in a vastly parallel manner. In other words, computation in CAs must be self organized, bottom up and driven by local interaction which are the features that all complex systems share due to universality. An example of such a computation given by Sipper is contour extraction of an image, highlighting how a parallel local computation can yield a global result. While contour extraction is computation it is not very general, and the initial conditions are set up in a rather artificial way, so one might ask if CAs can perform more complex computation. As it turns out the answer is yes, as even simple 1D CAs are *Turing complete*, enabling them to express any computation. This is of little practical use though, as Sipper puts it: “This is perhaps the quintessential example of a slow bullet train: embedding a sequential universal Turing machine within the highly parallel cellular-automaton model”, however it does show that more complex rules or models are not inherently more capable.

Order, Criticality and Chaos

Contour extraction and turing machines are artificially constructed computations, neither spontaneously occurring or particularly complex. A pioneer in the field, Stephen Wolfram, was interested not only in expressing computations with CAs, but in which conditions computation could spontaneously occur. He divided CAs into 4 classes based on their behavior, from least to most complex, as shown in fig [I - IV]. Class 1 and 2 quickly resolved to static or simple periodic behavior, which corresponds to an attractor landscape with many point and cyclic attractors with shallow basins. Class 3 CAs exhibited chaotic behavior where small differences in initial conditions caused the systems to diverge quickly into cycles so long they could effectively be considered to be in strange attractors. Class 4 was characterized by having very long transients, forming complex patterns of localized

²Singular: Automaton, Plural: Automata

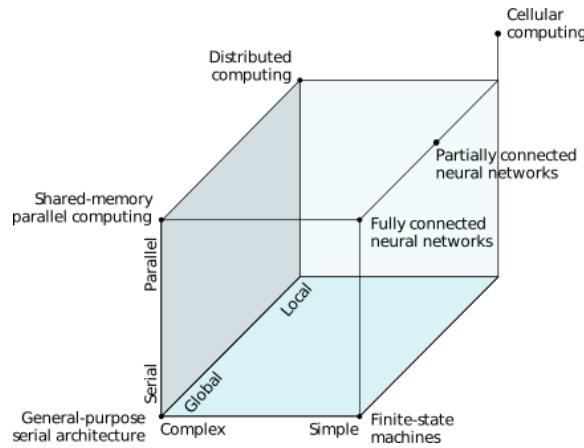


Figure 2.3: From Sipper paper.



Figure 2.4: From Bar Yam book.

structures corresponding to an attractor landscape with fewer attractors with very large basins. Wolfram correctly suspected that the fourth class was capable of computation, even of the universal³ variety. In Langton's pioneering paper *Computation on the Edge of Chaos* [3] Langton explores the space of possible transition tables for one dimensional cellular automata, linking the space of CA behavior to phase changes in material. Langton ascribed a parameter, λ , to each ruleset which described how evenly matched rules leading to the dead state was vs the live state. When all configurations lead to one type of cell the ruleset $\lambda = 0$, while the ruleset where the amount of states leading to dead and live are equal $\lambda = 1.0$ Langton found that the λ parameter modeled how rulesets and wolframs CA classes interact in a way analogous to how matter changes phase at different temperatures, shown in 2.8. Fig 2.6 is a phase diagram for water where the colored lines shows first or-

³Not to be confused with universality of complex systems

der phase changes, such as freezing below 0 degrees at ocean-level pressure, analogous to how CAs go from fixed to periodic in 2.8. More interesting, the figure also shows a *critical point* where matter undergoes a *second order phase change* in which it exhibits properties of several phases at once. In the critical phase matter forms structures that are self-similar, a property known as *scale invariance*, which causes classical assumption about “smoothness” of behavior at lower scales that allows simplifying the lower scale behavior break down. This is a recurring theme for systems, at the *edge of chaos* there is a critical point where the forces of feedback and dampening are closely matched, allowing a system to be responsive and dynamic without disintegrating into chaos.

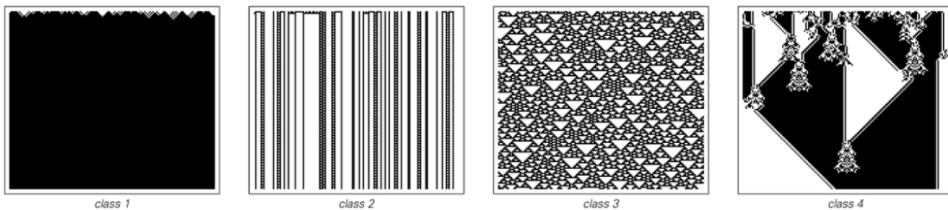


Figure 2.5: Source: A new kind of science.

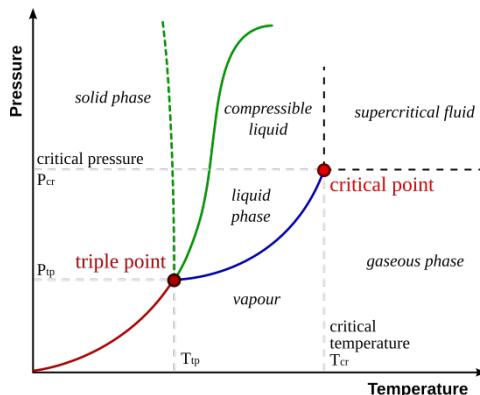


Figure 2.6: Source: wikimedia commons

2.2 Material Computing

Modeling neurons as cellular automata gives an idea of how neurons organize themselves into computationally capable networks and how this computation may look like, but the question of how to make sense of these dynamics and utilize them. Before tackling neurons it is necessary to extend our model to the physical realm and develop a model that can be extended to perform computation. A very natural extension of the CA model is *material computing* in which the atoms that make up matter locally interact in similar ways to CAs. Two pioneers in this field were the british duo Pask and Beer which studied how

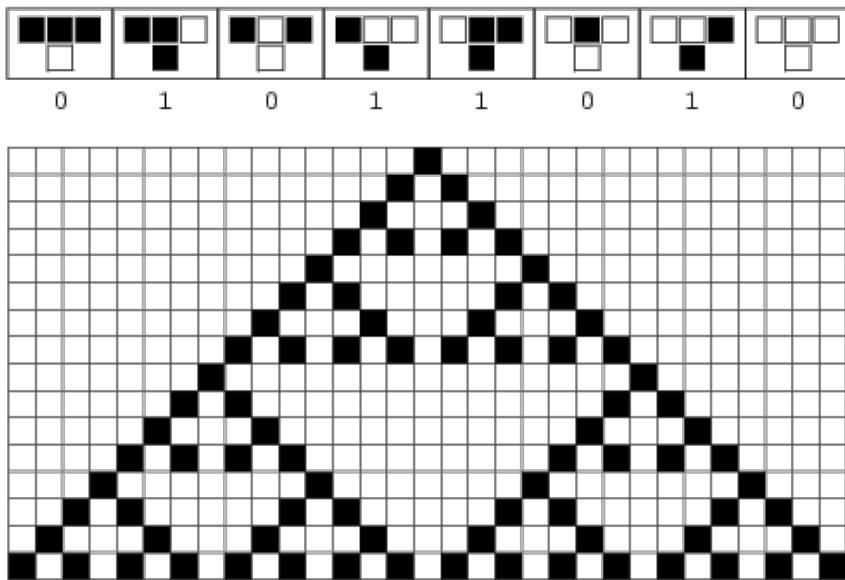


Figure 2.7: Should add a time axis to drive the point home that these are 1D

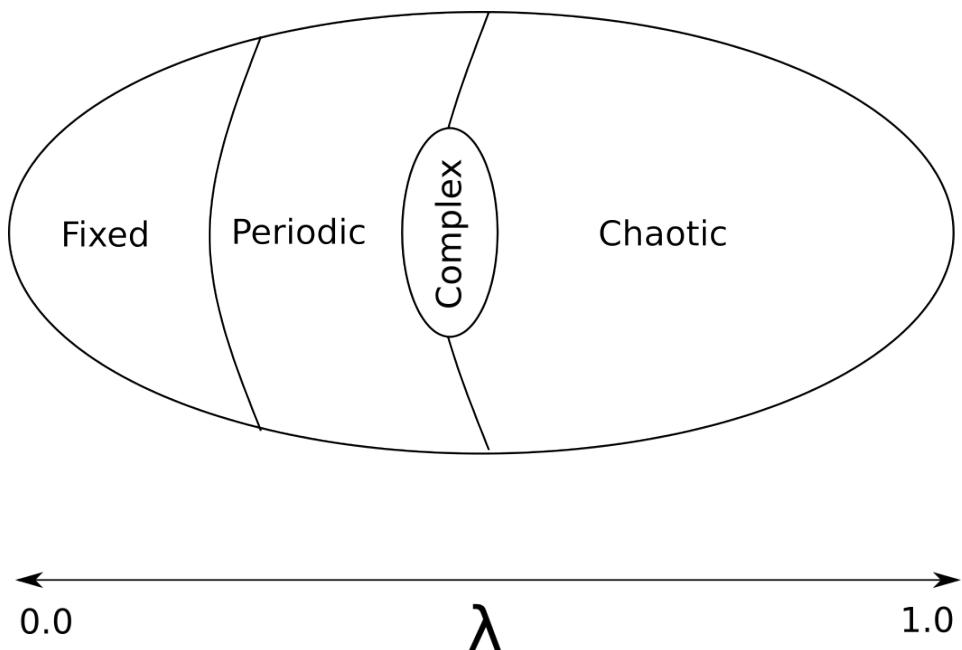


Figure 2.8: Langtons egg

unstructured matter could be used to perform computational tasks. In one experiment [cite ???] the duo used silver in an acidic solution which would form short-lived silver filaments when subjected to electric currents. By tuning the various knobs controlling parameters such as voltage this solution could be made to perform the task of tone discrimination, proving that unstructured matter could in fact perform computation. In Langton's CAs the goal was not to compute, but to explore what conditions were necessary for computation to occur, while Gordon and Pask sought to actually compute by observing how the matter reacted to sound and tuning the parameters in response. Tommaso Toffoli argues that “Nothing makes sense in computing except in the light of evolution” in his eponymous paper [cite TT], which separates Langton's compute capable CAs and the tone discrimination which had been evolved by repeatedly observing its performance and altering the parameters accordingly. A corollary of the turing completeness of cellular automatas is that CAs can be “tuned” in a rather heavy handed way to compute, but a more interesting conjecture is that a CA with a random seed will eventually exhibit local self-replicating behavior, spontaneously giving rise to what Toffoli would regard as computation, or possibly even as life.

Whereas the pioneers in material computing sought to perform specific tasks with their experiments more recent work has been made for general computation. One approach is *Evolution In Materio* where various materials have their parameters tuned algorithmically to perform a transformation between input and output. A recent effort is the NASCENCE project which has developed a physical device, “mecobo” that can host a variety of materials which can then be interacted with using an array of electrodes. By using the same basic loop as the early material computing researchers the Mecobo board has been used to create XOR logic gates using unstructured carbon nano-tubes [Cite Odd paper 1].

(DS)²

In the introduction Fernando's water bucket computer was briefly mentioned [cite bucket]. In their paper they refer to the bucket as a “liquid brain”⁴, performing calculations by observing the resulting wave pattern made by small motors. Unlike its biological counterpart the “liquid brain” does not evolve over time, in fact it has no structure to evolve at all, at least not at the chosen scale of observation. (DS)² is shorthand for dynamic systems with *dynamical structure*, and materials that exhibit this behavior can be used to study and model biological systems. In these material systems the *state space*, that is the set of reachable states, and the transition function between them evolve over time in tandem with the systems dynamics, both shaping them and being shaped by them. In [cite Odd Rune EiM (DS)2] Odd Rune shows that a mixture of table salt and water exhibits (DS)2 behavior using the mecobo board by showing that the system's response to perturbation evolves over time. Similarly, experiments using simulated artificial spin ice [Cite Johannes] shows that in a lattice of nano-magnets several underlying microstates map to the same macro-state as shown in fig ?? This highlights the importance of observation level when describing systems. One could choose to... Something something.

⁴Also in quotation marks in the original paper

2.3 Computing Neural Networks

Material computing provides both a theoretical framework and practical approach for modeling and interacting with the computational capabilities of living neural networks. The $(DS)^2$ viewpoint is a very natural fit when modeling how neurons organize into networks. The structure in a neural network continuously evolves as each neuron individually seeks out other neurons to forge new connections while letting old connections wither and die off. The dynamics of a neural network is the electro-chemical communication between neurons in the network. These signals are obviously shaped by the topology of the neural network, but they are also the driving force between the continued evolution of the networks topology in ways that we only have a rudimentary understanding of. Although both table salt and neural networks exhibit $(DS)^2$ behavior only the neurons do this on purpose. When viewing Pask's tone discrimination experient from a $(DS)^2$ viewpoint we can catch a glimpse of the same process that shaped the rules governing neural self organization. Although this was not the view of the experimenter, the tuning of parameters performed by Pask altered how silver filaments formed and decayed, which can be viewed as a rudimentary approximation of how neurons form connections. In Pask's experiment the observer tuning the knobs served the same role as evolution did in shaping neurons, and neither force were particularly concerned about exactly what the underlying structure of the material did as long as the behavior of the material performed as desired on a macro-scale.

The Computing Neuron

The neuron, or nerve cell, is the basic building block of both the human brain and the nerve system. The culmination of billions of years of evolution, the neuron is a vastly complex entity, featuring complex chemical pathways and gene regulatory networks. However, as discussed in the section on complexity, this behavior is necessary in order for evolution to function. As a corollary, most of the complexity of the neuron is not necessary for it to function, it is simply a byproduct of evolution, an implementation detail. It is the view of the author that the behavior of a neural network can be modeled with a much simpler neuron such as in artificial spiking neural networks.

Although the computation performed by neurons can be modeled with simpler models, the fact of the matter is that the cyborg project utilizes real neurons rather than a digital approximation, so a cursory introduction of the neuron is in order. While there exists a multitude of different neurons in the human body it is sufficient to consider a simplified model neuron, shown in fig [a figure of a neuron]. The three main parts are the body, or *Soma*, the *Dendritic network* and the *Axon*. The dendritic network acts as a receptor sensing electrical activity around the neuron, while the axon transmits electric pulses to neighboring cells. The connection between two neurons is called a *Synapse*. In addition to chemical signals neural networks communicate and regulate their behavior through chemical signals. These chemical signals, known as neurotransmitters, correlate strongly with electrical activity however, thus measuring the chemical gradients in neural networks is not a priority from a computational perspective.

Neural Dynamics

The “medium” of communication between neurons chosen as the observable macro state, that is the dynamics we are interested in measuring for the neural network are electrical bursts of activity, known as spikes. Figure [fig of waveform] shows recorded activity from an electrode. Each electrode measures the activity of the surrounding neurons rather than the state of a single neuron. Given that spikes propagate through the network in a cascading manner this level of observation is sufficient and more fine grained measurements is not a priority. Over time these dynamics evolve in tact with the structure of the network, as expected in the section on $(DS)^2$. Not only does the dynamics of the neural network change like the table salt experiment, it moves in the complexity space as shown in [fig neuro complexity graph] towards a critical phase as discussed in the section on cellular automata.

Artificial Neural Networks

In recent years ANNs⁵ (shorthand for artificial neural network) have revolutionized AI due to their ability to solve hard classification problems such as image recognition. The simplest, and most successful variation is the *feed forward* topology shown in 2.9 in which each layer only interacts with the layer directly in front and behind it. Shown in 2.10, the model for a single neuron consists of a summing function that calculates a weighted sum of incoming connections and a function that calculates the outbound activation of that neuron. More advanced models of neurons often retain state from their previous activation, one such variant being the spiking neuron which emulates the biological neuron more faithfully. Networks with neurons that retain state can have more interesting topologies where the connections between neurons can go in all directions. Cleary recurrent spiking neural networks should be more capable than feed forward networks, if nothing else for their ability to encode time-series as part of their state, yet the simple, non-spiking variant has proven much more successful. This gap stems from the relative simplicity of the fitness landscape of simpler network topologies which allows training algorithms to manually adjust the weights between layers iteratively until a good result is achieved. In back-coupled networks however the relation between cause and effect is far more obscure, thus algorithms which essentially assign blame and alter weights based on this are unable to cope with these topologies.

2.4 Reservoir Computing

Material computing has provided a tractable model of how the neuron computes and how to physically interface with them. However, still absent is a model that allows us to actually access the computational capabilities. In fact, from the field of material computing Susan Stepney extends the following warning “the biological substrate is extremely complex and complicated, having evolved over billions of years to exploit specific properties. In some sense, biological substrate is as far (or further!) removed from a primitive substrate as are our own designed abstract digital computational media.” As the heading indicates, the ace

⁵Or rather, the advent of hardware capable of efficiently running and training them.

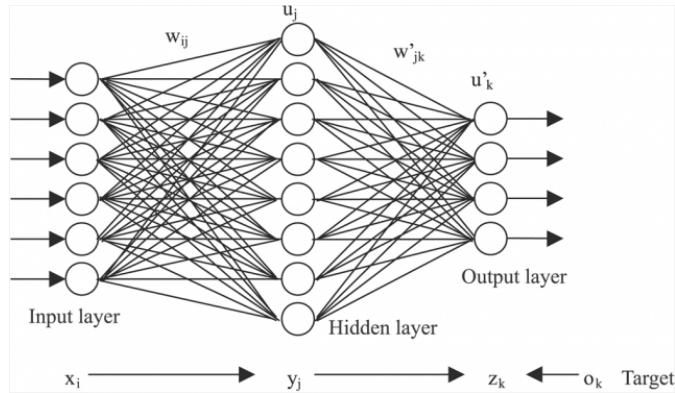


Figure 2.9: A feed forward ANN.

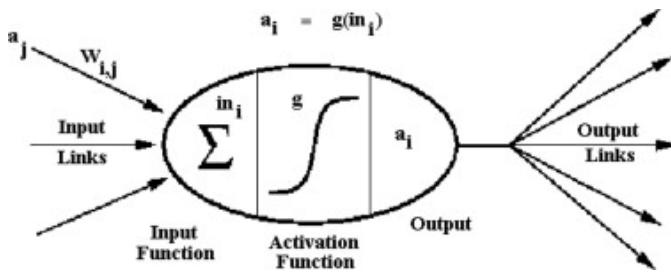


Figure 2.10: A feed forward ANN.

up our sleeve comes in form of *Reservoir Computing*, a technique that fittingly enough originated from the study of ANNs of the back-coupled variety[Cite Jaeger][Cite Maass]. In [Cite Jaeger] randomly connected neural networks are used to solve classification problems, forgoing training the network altogether. Instead the randomly generated network was used as a *reservoir* of dynamics that reacted to the input while a simple linear output function had to be trained. The principle behind reservoir computing is to employ a complex system as a reservoir which, as Schrauwen puts it [?] "... acts as a complex nonlinear dynamic filter that transforms the input signals using a high-dimensional temporal map, not unlike the operation of an explicit, temporal kernel function." Figure 2.11 shows this setup, in which input perturbs a reservoir and the resulting dynamics is classified using a linear filter. In order to explain why this works, schrauwen makes a comparison to the machine learning technique of source vector machines work, as shown in fig 2.12. The general idea behind an SVM is to expand the input into a higher dimensional "feature space", which in the figure is represented by the transition from 2D to 3D. Similarly, when perturbed by some initial condition the resulting dynamics from the reservoir can be interpreted as a feature space and a trained *readout layer* can interpret the resulting dynamics. Schrauwen points out two major differences between SVMs and RCs. First, SVMs only implicitly expands the input to high dimensional space in order to make the problem tractable, while reservoirs do not. Secondly, kernels are not capable of handling

temporal signals. The second difference is very important, it is what allows reservoirs to implicitly encode temporal signals in their dynamics, making reservoirs a natural fit for tasks such as speech recognition where each part of the input is context sensitive.

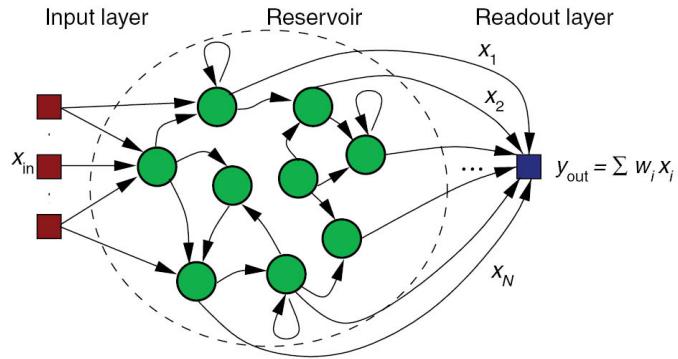


Figure 2.11: A feed forward ANN.

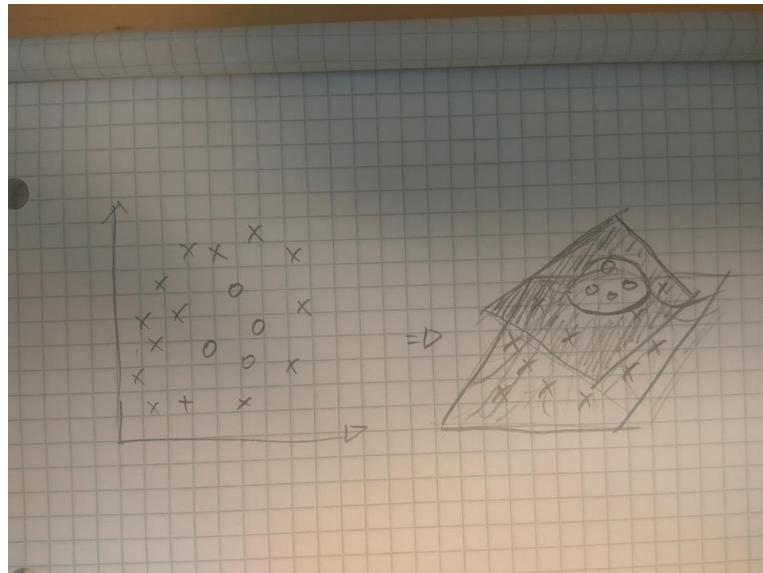


Figure 2.12: Should say “kernel function” or similar on top of the arrow

2.4.1 Linear and nonlinear output layers

TODO: Elaborate on the use of linear vs nonlin classifiers.

Chapter 3

Making Of A Cyborg

What is real? How do you define 'real'? If you're talking about what you can feel, what you can smell, what you can taste and see, then 'real' is simply electrical signals interpreted by your brain.

Morpheus to Neo - The Matrix

In order to get mad grant \$\$\$ we made a cyborg. Hello, here is some text without a meaning. This text should show what a printed text will look like at this place. If you read this text, you will get no information. Really? Is there no information? Is there a difference between this text and some nonsense like "Huardest gefburn"? Kjift – not at all! A blind text like this gives you information about the selected font, how the letters are written and an impression of the look. This text should contain all letters of the alphabet and it should be written in of the original language. There is no need for special content, but the length of words should match the language.

3.1 Concept

In figure 4.4 the conceptual cyborg is shown. This conceptual cyborg is comprised of three main components: An *MEA*, short for *Micro Electrode Array* in which a biological neural network is grown. A *neural interface*, allowing two-way communication between the neural network and the outside world. A robotic body, responding to movement commands and equipped with a sensor allowing it to perceive its environment. In this concept neural readouts are transformed into a Left and Right signal by a feed forward neural network, controlling the direction of the robot. Simultaneously data retrieved from the sensors of the robot are processed in a feedback processor and fed back to the neural network. The

conceptual cyborg is a closed loop system: The only input to the system is what the sensors perceive which the cyborg must act upon.

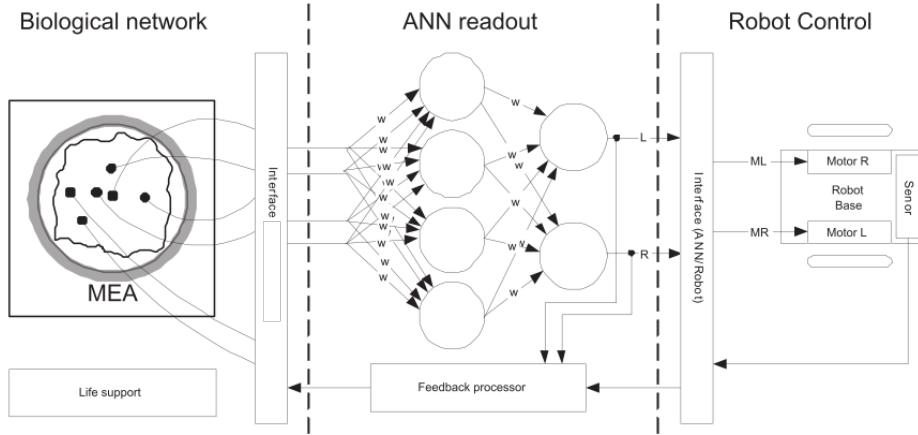


Figure 3.1: A simple conceptual cyborg.

3.2 Architecture

As shown in fig 3.2, the actual cyborg adds a lot of parts compared to the initial concept. This overview focuses on the main data loop, depicted as red and green lines, with red denoting the data pipeline from the reservoir to the reservoir computer, and green from reservoir computer to reservoir. This design expands on the initial concept, fleshing out the “interface” box in 4.4 and substituting ANN with the more general notions of reservoir output filter. The expanded architecture reveals a very important detail in the final design, which is that the data loop between the reservoir computer and the reservoir extends over network. This design choice allows the neural cultures to be safely stored in a laboratory, which has ramifications both in a practical and philosophical sense. By decoupling the reservoir and reservoir computer, the neurons are not restricted to a single robot. In fact, the robot does not have to be a physical robot, it can be fully virtual without the neurons noticing any difference. Additionally, the reservoir computer can interface with any reservoir as long as the interfaces are adhered to, allowing other reservoirs to be used in place of the neurons, both for testing purposes, and to compare the capabilities of different reservoirs.

3.3 Wetware

As opposed to hardware and software, the term wetware describes system components of biological origin, i.e “wet” components. The wetware of the cyborg is thus the neural networks which are being grown in MEAs at the department of neuroscience located at

St.Olavs research hospital. The MEAs are seeded with neural stem cells of either human or rat origin which then spontaneously form networks. At seeding there is no network at all, only a “soup” of dissociated neurons which over the course of several weeks start forming networks. The activity from these so-called pacemaker neurons can be seen in ?? . In the figure each cell in the grid corresponds to one of the electrodes as seen in ?? , however at this stage the monotonic spiking activity tends to be transient, starting and stopping randomly.

As shown in fig 3.4 a neural culture evolves over time, showing an increase in behavioral complexity. TODO: should say something more about this.

3.3.1 Neural Interface

The *Neural Interface* is the bridge between the biological and digital, responsible for providing a convenient API for reading out digital waveform data and inducing stimuli. Although not intended for use in close loop systems, the *MEA2100* system features the necessary hardware to implement a neural interface with the necessary functionality for a closed loop system. The *MEA2100* is built to conduct in-vitro experiments electrically active cell cultures such as neurons contained in micro electrode arrays by measuring and inducing voltages with its high precision electrodes. In addition to voltage and current the *MEA2100* system comes with a life support module, allowing neural cultures to survive for prolonged experiments for up to 30 minutes. Figure ?? [TODO: Figure missing] shows a physical overview of the components that make up the neural interface, which is also marked in 3.2.

Micro Electrode Array

Shown in ?? [TODO: Figure missing], an MEA containing a neural culture is slotted into the headstage for measurements. Not shown however is that the MEA can easily be removed, and is in fact one of many MEAs which are kept in an incubator when not experimented on. These MEAs feature 60 electrodes, which when accounting for the reference electrode provides 59 individual measurement and stimuli points, evenly spaced out in a grid. Each MEA is seeded with a neural culture, meaning that once seeded an MEA will be the host of a single culture, each capable of living for over a year, like the culture shown in ?? .

Headstage

Equipped with 60 high precision electrodes, shown in 3.5, the headstage can connect to an MEA when inserted into the measurement bay of the headstage. Thanks to the headstage each MEA is relatively expendable as they need only provide electrodes, leaving the measurement to the headstage.

Interface board

The interface board connects to up to two head-stages and is responsible for interfacing with the data acquisition computer, as well as auxiliary equipment such as temperature

controls. For this project, the most interesting feature of the IFB is a user programmable digital signal processor (DSP), which can access the raw stream of data from the headstage and issue stimuli. This DSP has access to the interfacing hardware in the IFB which is otherwise off limits and not reprogrammable for the user.

3.4 Software

Creating a cyborg is a massive undertaking, thus a quite extensive software suite has been created to make research feasible. As the system architecture overview 3.2 hints, the software is the main body of work performed for this thesis. Another important detail is the strong emphasis on the dataflow which is what connects the individual components. What goes on inside a component is less relevant compared to the transformation of data, visible from the outside. Finally, the figure shows a division between two systems, MEAME and SHODAN. MEAME runs close to the neural cultures, residing on the DSP and lab computer, and is specialized to interface with the MEA2100 system specifically. SHODAN on the other hand resides on the other side of the “network chasm”, and is closer to a framework, capable of handling any reservoir as long as the necessary transformations are implemented.

3.4.1 MEAME

An overview of MEAME is shown in 3.6, revealing that MEAME internally consists of two distinct subprojects exposed by a unified REST interface. While typically REST interfaces imply that a service is public, the MEAME REST interface is only intended to be used by SHODAN. Behind the REST interface there are two modules, data acquisition and DSP interface. The data acquisition module is responsible for configuring recording parameters such as samplerate and starting or stopping recordings, while the DSP interface provides a very thin layer of abstraction for writing to the DSP memory. Both these modules are built on top of a very thin API provided by the equipment vendor, making it clear that the closed loop cyborg system is pushing the equipment further than the vendor intended, for better or worse. In addition to the REST interface MEAME exposes raw TCP sockets, outputting the raw waveform data once the lab equipment has been configured. The format of the raw output is shown in figure 3.7, which must be demultiplexed by the receiver in order to separate data into individual electrode readouts, referred to as *channels*.

MEAME-DSP

Housed on the IFB, the DSP is not in use during normal operation and thus fully user programmable. Communication between MEAMEs DSP interface and the DSP itself is done with a primitive remote procedure call system, implemented using shared memory on the DSP. When SHODAN wishes to execute a DSP procedure this is done by submitting a list of memory writes via the REST interface, and then reading a special register that is incremented when the procedure call has been executed by the DSP, indicating that a new call can be executed, and that return data is valid and can be safely read. This rather thin API is admissible because of the fact that MEAME is only accessed by SHODAN

which internally translates procedure calls to memory writes and reads. The stim control module maintains a list of *stim groups* which contain a list of electrode numbers and a desired frequency. Consequentially, stimulating neurons is done by specifying which electrodes should be stimulated, and at which frequency, rather than explicitly sending a signal whenever stimuli should be applied. A visual representation is given in 3.9 showing two stimulus groups and the resulting stimulus applied. The stimulus group configuration also specifies a pattern for each group, sine and square respectively. These patterns can be uploaded by directly writing to memory from SHODAN, but for this project simple square waves are adequate.

3.4.2 SHODAN

As shown in 3.2, SHODAN is responsible for storage, filtering, generating perturbations and maintaining the core RC loop. During an experiment SHODAN can be divided into two parts, reservoir interfacing shown in 3.8, and the core RC loop shown in ???. The former can in turn be subdivided into input and output processing with regards to the main RC loop.

Input Processing

The MEAME comms interfacing module mirrors the corresponding interface exposed by MEAME, consuming raw waveform data over a TCP connection and configuring experiment parameters and issuing stimuli requests using a HTTP client. The channel demultiplexer splits the raw datastream from the TCP sockets into individual channels as described in 3.9 before entering the filtering module. The filtering module is responsible for translating the raw waveforms to a format that can be utilized by the main RC loop. The currently implemented filter does a basic “spike detection” transform as shown in ???. When a certain voltage threshold is reached a spike is recorded and the detector is disabled for a short cooldown period before it can detect the next spike. The spikes are then input into a moving average filter, smoothing out the staccato spike/no spike filter output to the amount of spikes that has happened between t and $t - \Delta t$ where Δt is on the order of 100ms to 1000ms. The final input to the core RC loop is a vector of values, one for each channel that can be periodically sampled. As the data crosses the boundary to the core RC loop, its shape reveals little facts about its origin. For any other reservoir it would be just as natural to translate the output dynamics to a vector of scalars, thus for any reservoir the filter module is what translates from specific to generic.

Output Processing

Being the dual of the Input processor, the output processor does the same thing in reverse, transforming data from generic to specific. In the figure 3.8 the output of the main RC loop is shown as a list of distances, but that is merely the interpretation of data, the underlying shape is just the same as the output of the filter from the input processor: a vector of scalars. Since the goal of the project is to create a cyborg the focus is naturally on robotics, virtual or not, but the main RC loop can just as well handle other tasks such as classifying images, and the resulting output would still be expressed as a vector of scalars. Figure 3.10 shows

this duality between the input and output processor: Just as the input processor masks the identity of the reservoir from the core RC loop, the output processor masks the identity of the task being performed!

The datapath of the output processor is rather similar to the input processor with the perturbation transform module serving the same purpose as the filter in reverse, turning scalars into frequency ranges. In the figure an additional module is shown, transforming a stim request into a set of memory writes for the DSP as described in the MEAME section, but this is just a practical matter added for completeness.

Core RC Loop

In the conceptual cyborg 4.4 all of the software described previously can be fit into the interface between the ANN and MEA. The core of the reservoir computing system is the ANN readout and the robot control module, but the concept leaves out a very important piece of the puzzle: How can the correct ANN readout layer be chosen? The Core RC loop consists of three components, the robot control, the RC filter, and the reservoir performance evaluator. Just as the robot control does not have to be an actual robot, the RC filter can employ any filter, it does not need to be an ANN. The last module is what separates SHODAN from a classical reservoir computer. In classical reservoir computing the reservoirs themselves are fairly static, exhibiting little changes in behaviour over time. A reservoir computer using a random boolean network can store an RBN in memory and access this data to reconstruct the experiment at any time. This simplifies the task of training the RC output layer since the reservoir will not change between runs. Unlike digital RBNs, neural cultures have no reset button, their behaviour changes both long term and short term. Long term, a cluster of neurons might “emigrate” away from an electrode, new connections form and old connections wither. This process happens over days and weeks, but change also happens on a short basis, both due to natural fluctuation, and because of the fact that no two experiments are truly independent, there are no ways to make neurons forget and revert to previous behaviour. The solution to this problem is to back link the robot control to the RC filter via a performance evaluation module, creating a feedback loop capable of coping with the moving target presented by a neural culture.

3.4.3 Storage And User Interface

This section will feature the UI and a description of the database system and the sort of offline analysis that can be done. It’s not that interesting, but it is significant.

3.5 Current State

MEAME and SHODAN are both still in development and are far from complete. All functionality described in this chapter has been implemented and is in a running state, but the stimuli control module that is responsible for applying voltage to the electrodes does not work correctly. This means that the closed loop system does not work at all until the issue can be fixed. This frustrating issue reflects how the closed loop system goes far

beyond the intended use for the lab equipment, which is marred by poor software and non-sensical documentation, meaning development is often a process of reverse-engineering. While this flaw means the system does not fulfill its purpose, it's enough to verify that the remaining modules work. Recordings can be made and played back, which means that MEAME/SHODAN is already able to provide a better storage solution than the vendor provided solution.

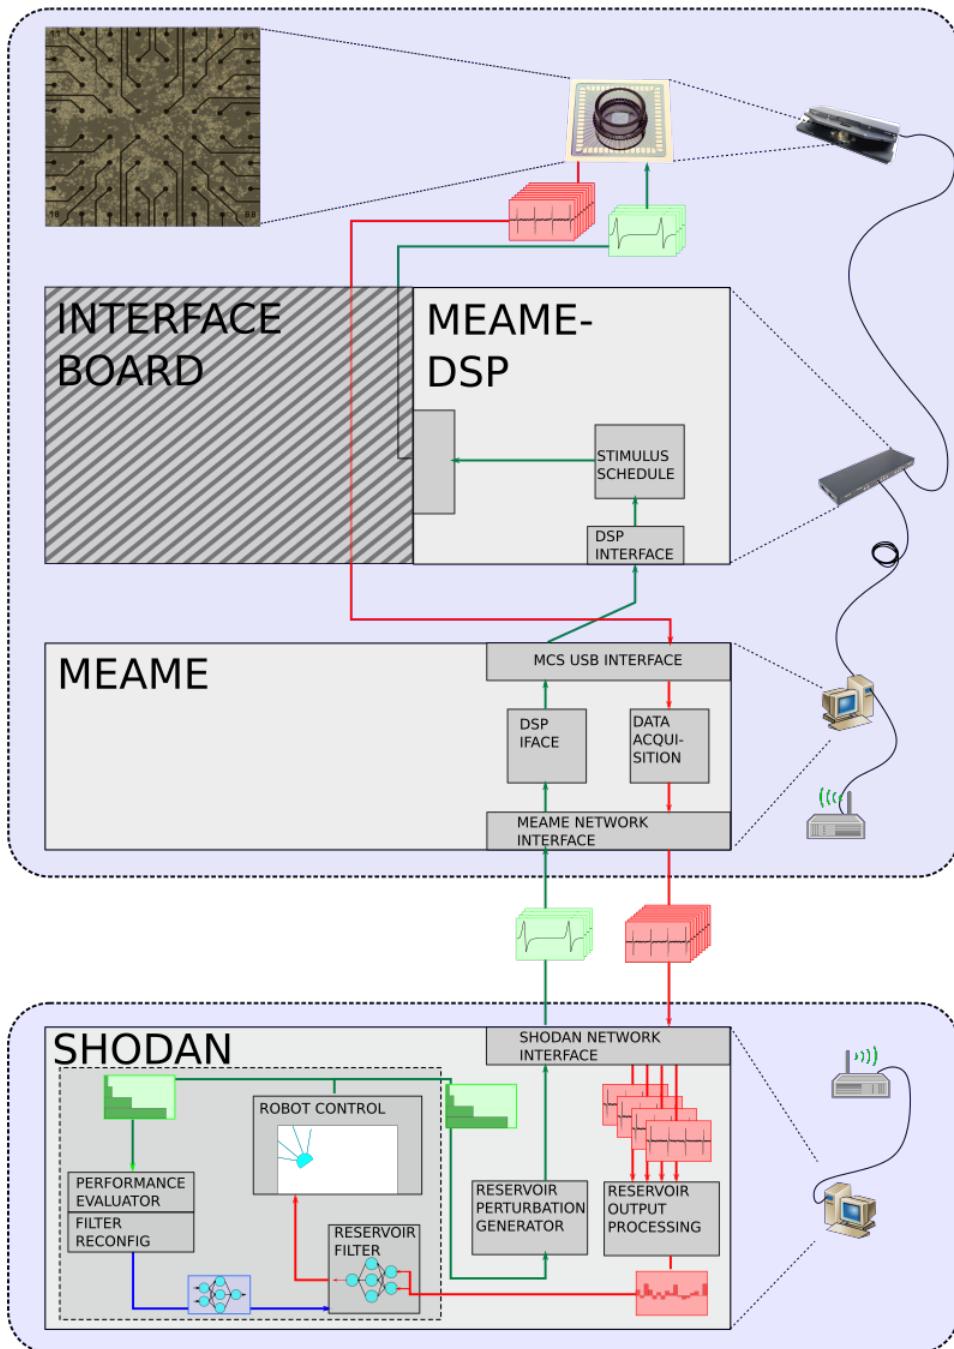


Figure 3.2: An overview of all software written for the cyborg project. The dataflow is color coded with red denoting data going from neurons to robot, green from robot sensors to neurons. The ble line indicates configuration data, as the system is capable of reconfiguring itself during an experiment. The two main systems are physically separated, connected only through network.

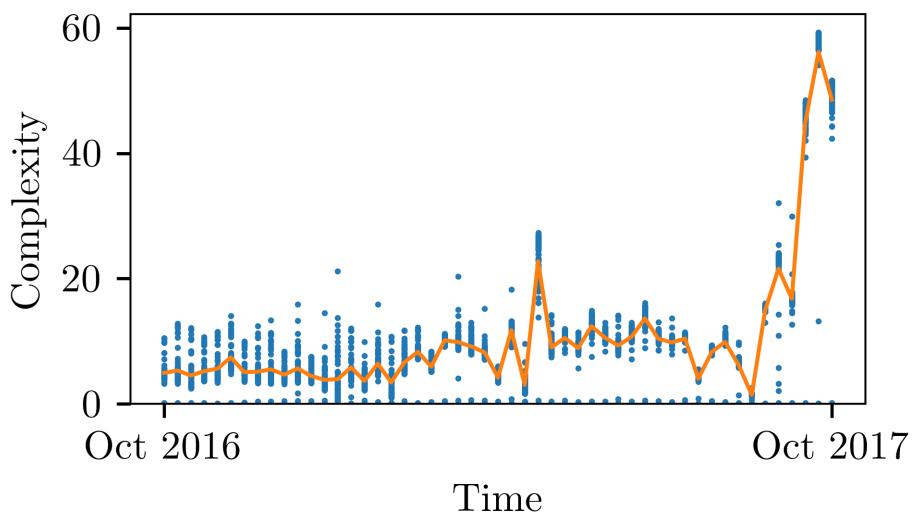


Figure 3.3: Approximated Kolmogorov complexity of the firing patterns over a one year period.

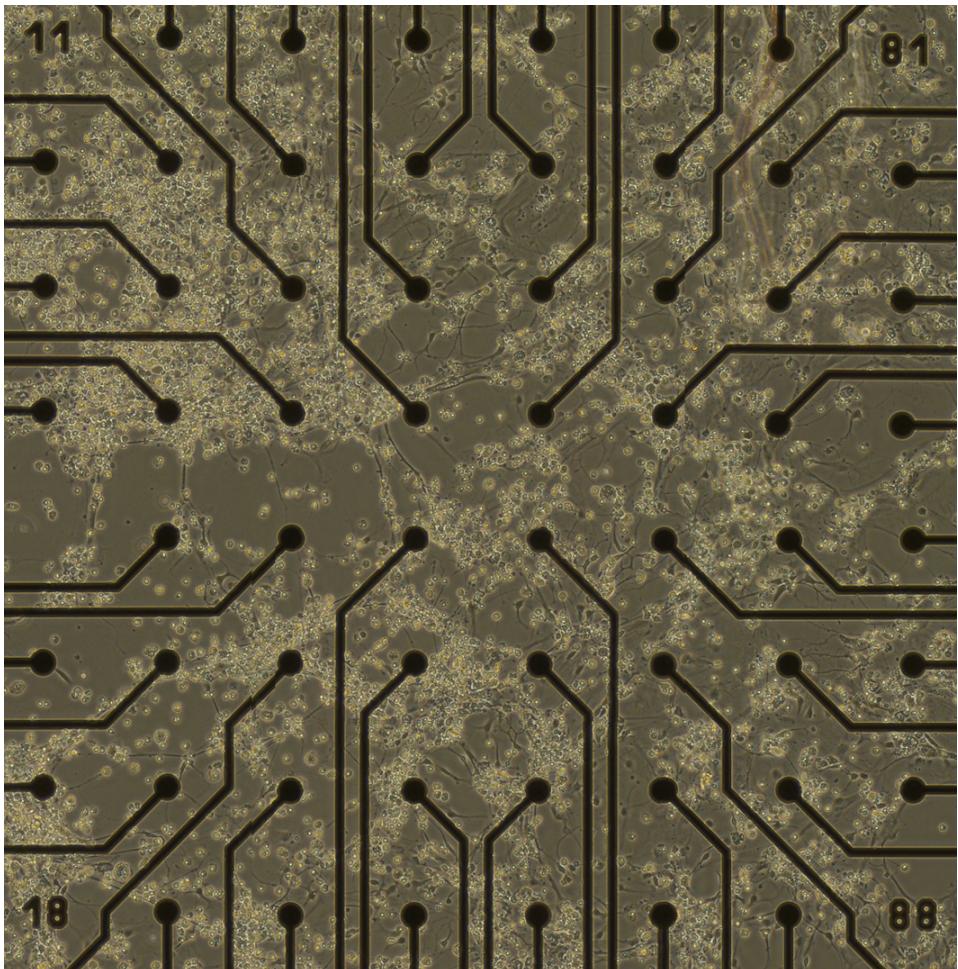


Figure 3.4: A look through the microscope, showing a culture nicknamed “Frank”, a play on frankenstein’s monster due to its tendency to develop *Organoids*, small proto-organs with structures similar to eyes and other sensor organs. The black lines are electrodes that can be used to record and apply stimuli.

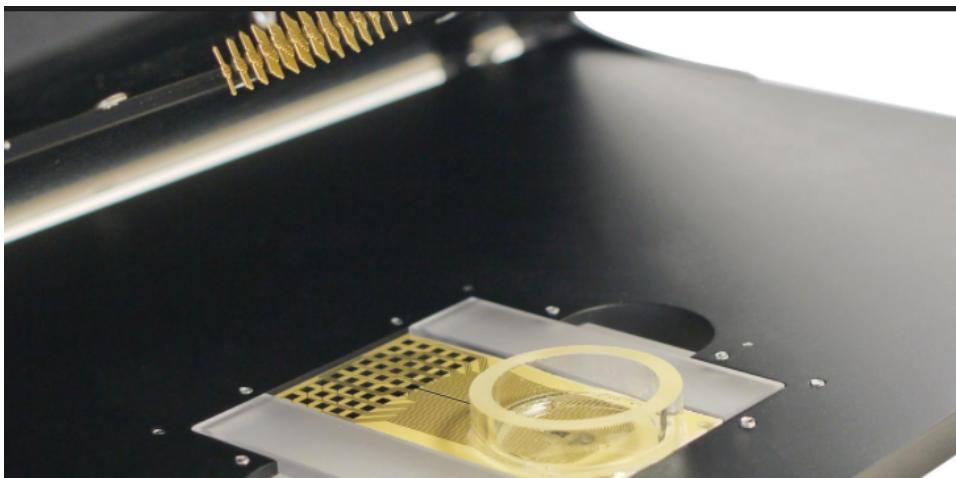


Figure 3.5: A placeholder for an open headstage showing the electrodes.

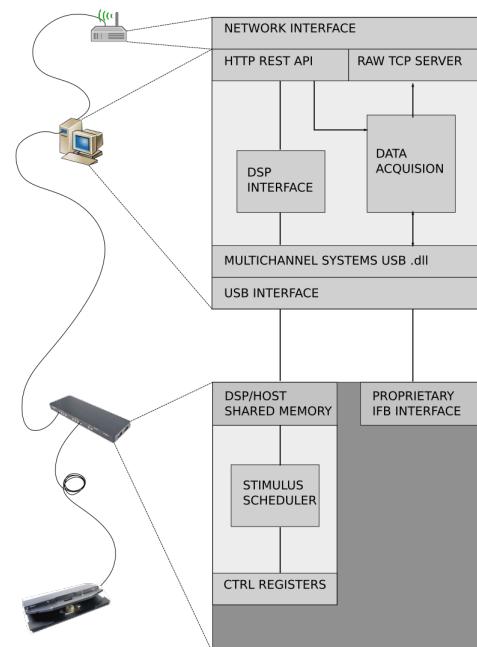


Figure 3.6: Rough sketch. nice MEAME TODO: Flesh out the DSP box with RPC handler a la MCS USB .dll

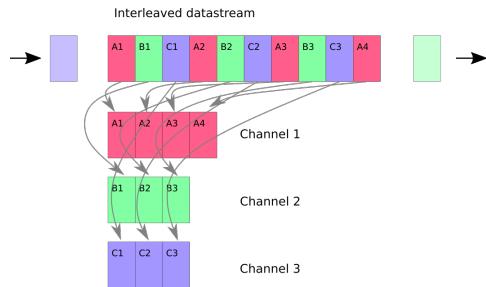


Figure 3.7: The interleaved datastream can be broken down into its individual channels. Here a datastream with only 3 streams are shown, whereas the stream is normally comprised of data from 60 channels.

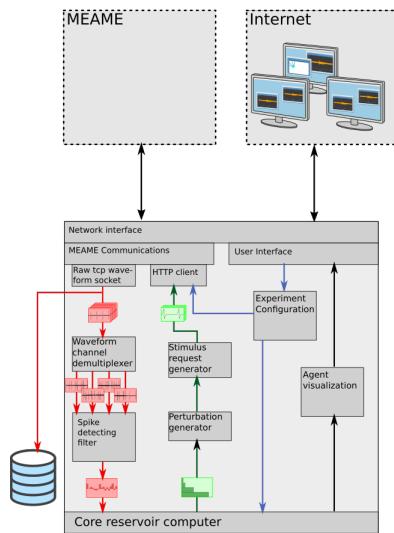


Figure 3.8: Rough sketch. Some wavez

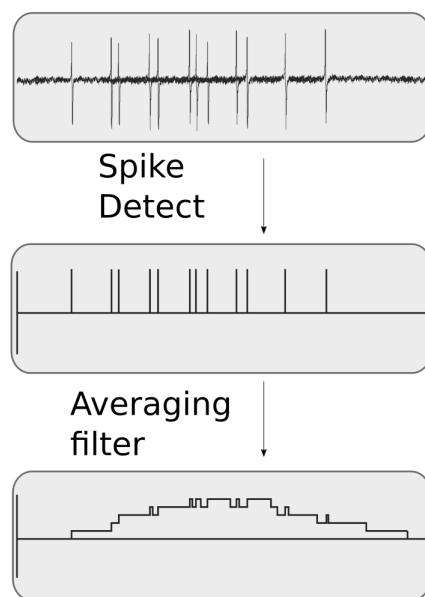


Figure 3.9: Rough sketch. Some wavez

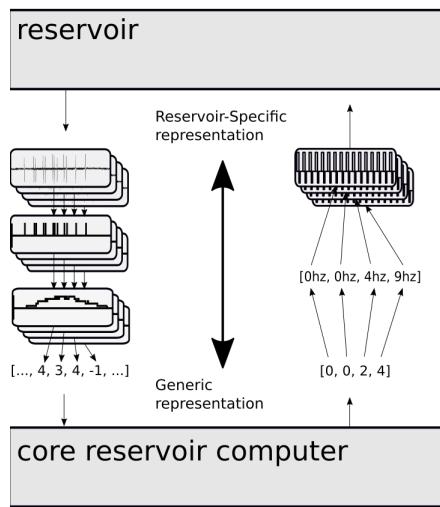


Figure 3.10: Rough sketch.

Chapter 4

Experimental System and Setup

Because of the issues with stimuli generation the experiment setup has not been employed on live neural tissue. As a consequence, the current experimental setup is designed to be easy to modify whenever testing becomes possible and to verify functionality of the rest of the system. Consequentially, the focus on this section is the experimental capabilities provided by the system and the parameter space that can be explored. When leaving out all implementation details the experimental setup provided by SHODAN and MEAME becomes much simpler, as shown in 4.1. The figure highlights a point mentioned in the previous section, namely that the closed loop system actually consists of two data-loops running in parallel. This is a departure from the classical reservoir computer model where modifications are made between runs. In the classic model experiments are run, performance is evaluated and the readout layer is altered in some way. This mode of operation makes sense when the reservoir can be reset to an initial state, but as discussed this is not possible when dealing with neural tissue which evolve both short term and long term. The experiment setup from 4.1 can be broken up into three parts: The data loop, the core reservoir computer and the reconfiguration loop. The three following section gives a brief description of purpose and which parameters can be altered. To give a better overview these parameters are then summarized in the subsequent section.

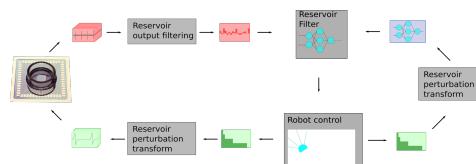


Figure 4.1: Rough sketch. Some wavez

4.1 Data Loop

A closer look at the data loop is shown in figure 4.2. The two “knobs” that can be turned in this part of the system is the reservoir input filter and the perturbation transform. The reservoir output filtering module, as discussed in the previous chapter, perform a spike detection which is then used as input to an averaging filter. While the spike detection is considered to be fixed, the averaging filter is not. In the current implementation it is a simple averaging filter, but it can easily be changed to use a different model, such as linear or exponential decay if the averaging filter is deemed insufficient. The second module is the reservoir perturbation transform which transforms a set of distances perceived by the robot into stimulus frequencies. The perturbation model has three parameters: A list of electrodes to receive the stimulus signal, the period between applying stimuli, and the stimuli signal itself. Much like the choice of using a simple averaging filter, the default settings for the perturbation transform favors simplicity. For each sensor on the robot a single electrode is chosen and is generally considered outside the interesting parameter space. Similarly a simple waveform is chosen, either a sine wave or a square wave. These samples are fixed, increasing the period means there will be a longer period between each sample, but the samples themselves remain fixed. Again, for simplicity, a simple square wave is the default choice, thus the only parameter that is not considered fixed is the transform between distance and period. The chosen default for this transform is an exponential decay, when the sensor is facing a wall directly the square stimuli frequency is set to 30hz with an exponential dropoff to 1/3hz when the sensor perceives a wall at the maximum distance.

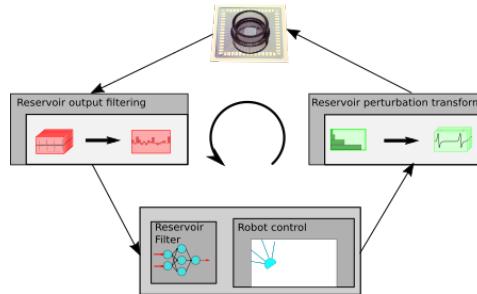


Figure 4.2: Rough sketch. Some wavez

4.2 Core Reservoir Computer

As discussed in the previous chapter, the core reservoir computer is the ideal reservoir readout layer and the perturbation generator while the remaining functionality is abstracted away in the interface and feedback processor box. The core reservoir computer consists of two parts, the robot control and the reservoir readout layer.

Readout Layer

The chosen readout layer is a simple feed forward neural network. As mentioned in the background, the choice of using a network with multiple layers as indicated in the figure means that a non-linear problem may be solved entirely by the readout layer. For the proof of concept the topology of the network is not fixed to a topology, thus it is up to the experimenter whether a network with a single layer or one with multiple layers should be used. In fact, there is no strict requirement that the readout layer must be an artificial neural network at all. While it might seem a wise choice to use a readout layer that has similarities with the reservoir itself, this line of thinking ignores the “anonymization” that happens between reservoir and reservoir computer, as shown in the fig 3.10 in the previous chapter. The final word on the choice of readout layer is that the artificial neural network approach has the following benefits: It can easily be extended from a linear to nonlinear classifier, and it is easy to modify during an experiment.

Robot Control

The robot control is a simulator that controls an agent with 4 “eyes” in a box world as shown in 4.3. As with all other parts of the system variables such as the amount of eyes, sight range, the function between the output of the reservoir filter and robot movement, max turn rate and max speed can be changed on a per-experiment basis. The goal of the “game” is for the robot to not collide with walls. Since the best way to avoid this is to simply run in circles, the performance is evaluated by how close the agent came to crashing in a series of trials where the agent is faced towards the wall at different angles as shown in fig [TODO]. Since the actions of the robot are decided by the output of the readout layer it the action of the agent indirectly depends on the evaluation function since it is this evaluation that shapes the next readout layer.

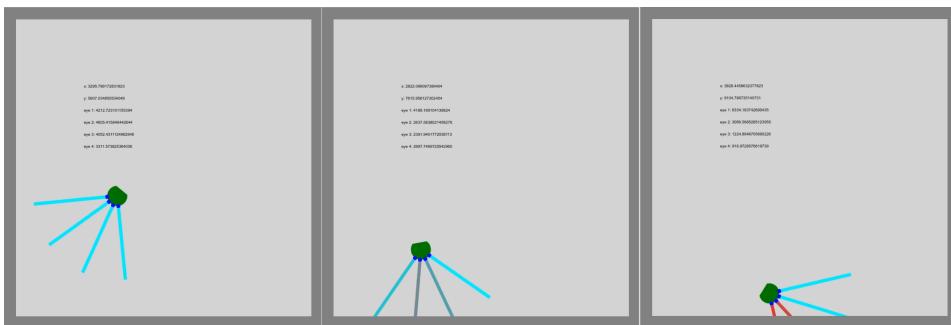


Figure 4.3: The cyborg doing its thing

4.3 Reconfiguration Loop

The reconfiguration loop, shown in 4.5 is the dataflow responsible for the on-line reconfiguration of the reservoir readout layer. On experiment start the reservoir filter reconfigurator

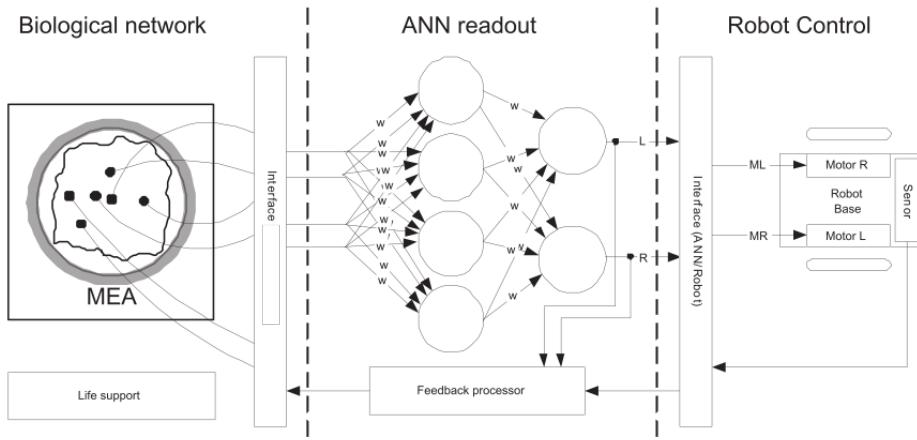


Figure 4.4: A simple conceptual cyborg.

supplies the readout layer module with a set of weights for the artificial neural network and records the resulting behavior. Each attempt is scored according to a scoring function as described in the previous section, and from this the reservoir filter reconfigurator generates new readout layers. The chosen method of readout layer generation is a *genetic algorithm* because of its ease of implementation and relative lack of bias.

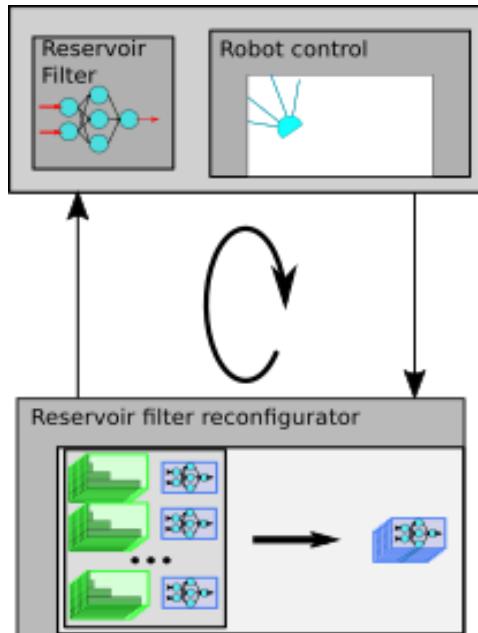


Figure 4.5: Rough sketch. Some wavez

4.4 Parameter Space

The first distinction in the parameter space is uhh...

Chapter 5

Conclusion And Further Work

The future will be better tomorrow.

Vice president Dan Quayle (R)

The goal of this thesis has been to show produce a working proof of concept for a neuro-digital hybrid. While this has not been achieved the necessary scaffolding is in place, and the system is capable of performing to specification in all cases save apart from application of stimuli. In addition to the stated goal a relatively sophisticated storage solution has been implemented, and a web-based user interface is being worked on, which will allow automation of the experiment workflow once it is further developed. While far from complete the system is already capable of performing most tasks currently handled by the vendor supplied software. While replacing this software is not a project goal, it gives an idea of the effort necessary to make the lab equipment behave.

5.1 Further Work

While the lack of results make it hard to draw a scientifically valuable conclusion it at least makes the further work section easy to write. The obvious first task is to figure out what is causing stimuli to not work which is simply a matter of time. TODO

Bibliography

- [1] , ???? Marvin Minsky - Wikiquote.
URL https://en.wikiquote.org/wiki/Marvin_Minsky
- [2] Fernando, C., Sojakka, S., Sep. 2003. Pattern Recognition in a Bucket. In: Advances in Artificial Life. Lecture Notes in Computer Science. Springer, Berlin, Heidelberg, pp. 588–597.
URL https://link.springer.com/chapter/10.1007/978-3-540-39432-7_63
- [3] Langton, C. G., Jun. 1990. Computation at the edge of chaos: Phase transitions and emergent computation. *Physica D: Nonlinear Phenomena* 42 (1), 12–37.
URL <http://www.sciencedirect.com/science/article/pii/016727899090064V>
- [4] Vardi, M. Y., ???? Artificial Intelligence: Past and Future.
URL <https://cacm.acm.org/magazines/2012/1/144824-artificial-intelligence-past-and-future/fulltext>

Appendix

My appendix was surgically removed. Take that, evolution.